Omniscient AI

The Sovereign AI Appliance — Sales Booklet

Your AI. Your data. Your premises. Governed by Australian law.

How to use this booklet

This is the full sales narrative for the Omniscient AI offering. The 1-pager, the 5-page visual pamphlet, and the enterprise SLA are all derived from it. Product name for the appliance is **Omniscient Core** (placeholder — change once finalised).

1. The problem with AI today

Every enterprise knows it needs AI. Almost none can adopt it without swallowing four uncomfortable compromises:

1. **It never stops costing.** Per-seat and per-token pricing means the bill grows forever, in lockstep with how useful the AI becomes. Success is punished.
2. **Your data leaves the building.** Prompts, documents, and customer records are shipped to foreign clouds you don't control, retained on terms you didn't write, and — in many cases — used to train the next model.
3. **The guardrails aren't yours.** Foreign models, tuned by foreign trust-and-safety teams, refuse legitimate, specialist work — medical, legal, mining, defence, agronomy, engineering — because a generic policy doesn't understand your industry.
4. **Nobody's accountable to Australian law.** Offshore vendors don't build for the Privacy Act, the emerging Australian AI guardrails, or your sector's obligations. When the regulator asks where your data is and how the decision was made, "a US data centre, we're not sure" is not an answer.

Mid-market enterprises feel this hardest: big enough to have real compliance exposure and real volumes of sensitive data, not big enough to build a sovereign AI platform in-house.

2. The Omniscient answer

A complete, private AI workforce in a box — installed in your office, owned by you, governed by Australian law, with no data ever leaving the building.

Omniscient Core is a hyper-secure physical appliance, installed on your premises and handed over running. Inside it: a hardened operating system, local open-weight AI models, a fleet of pre-built agents and workflows, and an open-source equivalent of nearly every piece of business software — all pre-integrated, so the AI doesn't just *chat*, it *acts* across your stack.

You command intelligence you physically possess and control. Not intelligence you rent over the wire while someone else keeps your data.

3. How we solve each compromise

Your compromise today	The Omniscient Core
Cost that never stops	A fixed 3-year lease (fully R&D-claimable) + local inference. The marginal cost of a query is electricity. Unlimited seats. Unlimited agents.
Your data leaves	On-premises, single-tenant, air-gappable. No inbound internet path. Data never leaves unless a named human approves a specific, logged exception.
Foreign guardrails	Customer-owned, customer-tuned open-weight models. Guardrails set to *your* business, not a foreign safety desk. Locally fine-tunable on your own knowledge.
No AU accountability	Onshore by physics. Designed to support the Privacy Act and Australian AI guardrails, with an audit ledger an auditor can inspect — and kept current by subscription as the law evolves.

4. What's actually in the box

A layered, sealed appliance:

- **Secure hardware** — a **six-bay enterprise chassis**, mini-rack or rack-mountable, with the **latest AMD or Intel multi-core processors**, **up to 256 GB of RAM**, **up to 200 TB of encrypted storage**, and a **GPU sized to your team**. TPM, full-disk encryption, tamper-evident chassis, and an optional physical air-gap switch.
- **Hardened OS** — immutable, signed, self-healing. Nothing to drift, nothing to break.
- **Local AI runtime** — open-weight models running on the box's own GPUs. Fast, private, unlimited. Optional, *consented* burst to a frontier model for the hardest tasks.
- **The open-source software estate** — preloaded and **expandable across every activity you could want**: CRM, content-creation tools, accounting, document editors, spreadsheets, business intelligence, team comms, file storage, automation, and a knowledge base — the open-source equivalents of the tools you use today, all yours, all AI-wired.
- **Agents & workflows** — pre-built AI agents configured to your departments during onboarding. They draft, summarise, analyse, and — with approval — act inside your software.
- **Governance & evidence** — an append-only, tamper-evident audit ledger that records everything the AI does, so every privacy and compliance claim is provable, not promised.

No compromise on parity or integration. The open-source estate is chosen and tuned to match the tools you run now — no performance penalty for going sovereign — and the AI is wired across all of it so workflows span packages seamlessly. Where your business depends on a proprietary or outside service — a bank feed, a government portal, an industry API — **we ensure interoperability**: the box integrates with it rather than forcing you to abandon it. Sovereign doesn't mean isolated; it means *you* decide what connects, and every connection runs through the same consented, logged egress controls.

4A. The Operating System — what your team actually uses

The hardware is why your CISO says yes. The **operating system is why your staff love it** — it's what they open every morning. And it's not a concept: it's productised from a real, working build.

Everything you use in Claude is included, on your box. The full library of slash commands, skills, and named specialist agents — research, document authoring, code review, finance and accounting workflows, marketing, legal, design, data analysis, and hundreds of domain skills — ships pre-installed and runs on your local hardware. No separate subscription, no per-seat AI licence. Staff ask in plain language; an "automatic agent selection" router sends the request to the right local model — or, with consent, bursts to a frontier model for the hardest tasks.

Multi-stage processes become one-button tasks. The home screen is calm: a few large work-group tiles, not a wall of dashboards. Click one and that area's workflow automations appear as buttons. Each button is a *saved orchestration* — a sequence of agent and tool steps that would otherwise take a dozen manual actions — collapsed into a single tap. A non-technical staff member runs a sophisticated multi-agent workflow with one click; an admin can open that same workflow as an editable, inspectable node graph.

Your daily operations run themselves — on the GPU you already own. Recurring work is analysed, scheduled, and run on the local GPU overnight, then presented for review:

- **Email triage** — inbox sorted, spam blocked, edge cases queued, replies drafted.
- **Social & marketing content** — drafts generated per brand, staged for one-tap approval (nothing publishes without sign-off).
- **Accounts payable & receivable** — bills and invoices analysed, AP runs prepared, receivables chased, reconciliations drafted — ready for human authorisation.

The GPU is never idle: when no scheduled job is running, it backfills with research and opportunity work, so you get full utilisation of hardware you've already bought. Each run leaves an audit-ledger entry, and staff see — per day — what every agent has done, what's queued, and prior outputs.

Updates that explain themselves — zero configuration. When Omniscient ships a new capability, the box presents a plain-language explainer and a single choice: *Yes, add it to my workflows* or *No, not now*. On yes, it wires itself into the right work group automatically — no setup, no config, no consultant. The platform compounds in value over time while adoption stays your informed, one-tap decision.

4B. The Second Brain — your company's knowledge graph

This is the single biggest difference between Omniscient and any AI you can buy today. Conventional AI is *stateless* — it forgets every session, and what your people know stays locked in their heads. The Omniscient second brain is the opposite: a complete, enduring, graphical memory of the business that grows with every day it operates. Every company file and data source — documents, emails, contracts, designs, records, prior agent outputs — is ingested and **graphed into a machine-readable knowledge network**: not a folder of files, but a living web of entities and relationships (people, projects, customers, products, obligations, decisions) the AI can traverse.

- **Split by department.** Finance sees the financial graph, Legal sees legal, R&D sees R&D — under the same access boundaries as the rest of the box. Each team gets its own slice of the second brain.
- **It saves you money.** Because the AI has complete, pre-organised context, it doesn't need long documents re-pasted or the business re-explained on every request — it queries the graph for exactly the relevant entities. Fewer tokens per task means faster, cheaper inference (and far smaller, cheaper burst payloads), *and* better answers, because the model reasons over the whole relevant web of company knowledge.
- **You can watch it grow.** The graph is rendered as a live, interactive 3D visualisation — orbit, zoom, query it by voice or chat. As work happens, new nodes and connections appear. Teams literally **see their proprietary knowledge base grow alongside business activity.**

It's a flywheel: more work → a smarter second brain → cheaper, better automation → more work done. The box gets more valuable the longer it runs — and that accrued institutional memory is an **asset you own**, on your premises, in open formats. Not a vendor's data moat. Yours.

4C. The Employee Brain — capability, continuity, and zero-loss handover

The second brain resolves all the way down to the **individual**. As each person works, their slice of the graph becomes a living *employee brain* — and that delivers three things no cloud AI can, because no cloud AI keeps the data on your premises:

- **An honest measure of AI adoption.** You can finally see *how* each person actually uses the tools — what they automate, where they build on and share agents, where they're getting leverage and where they're stuck. The missing measurement layer for any AI rollout: real adoption per role, training needs surfaced, power users identified, ROI per seat made visible.
- **A complete, live capture of the role.** Active clients, open tasks, commitments, decisions, and current progress across all work — captured continuously as a byproduct of doing the job, not a status report someone has to write. The role's true state lives in the graph, not in one person's head.
- **Handover without errors or gaps.** When someone goes on leave, changes roles, or leaves, the successor inherits the *whole picture* — every client, every open thread, every commitment, every piece of context. Nothing walks out the door; no client falls through a crack. The classic risk of losing institutional knowledge to staff turnover — the "bus factor" — is structurally eliminated, and onboarding a replacement goes from months of rediscovery to a guided walk through an existing brain. **With Omniscient in place, the painful handover — and the loss of institutional knowledge when an employee leaves — is a thing of the past. The knowledge is all there, and it is handed to the next person with full context and every detail.**

> **Said plainly:** this is a *continuity and capability* tool, not a surveillance tool. It's captured by working through the box, governed by the same access controls and audit ledger as everything else, and deployed transparently to staff in line with Australian workplace and privacy law — covered in your compliance pack. The benefit is real; we give the answer to the HR/legal question rather than dodging it.

4D. Self-improving loops — the work gets cheaper and better every run

The second thing that sets Omniscient apart is that it *improves itself*. Conventional AI does the same task the same way every time, and waits to be told what to do. Omniscient runs your repetitive work and workflows as **self-improving loops**:

- After every run, it **assesses how it performed** — what worked, what it learned, where it can do better.
- It continually **re-evaluates** how it goes about the task: the resources, knowledge and pathways it uses.
- It **re-optimises** accordingly — constantly, run after run.

The result compounds: **each run uses fewer tokens, each task gets faster, and each project's outputs improve.** Cost falls and quality rises the longer the system runs — the opposite of SaaS AI, where the bill only grows with use.

This reflects a different operating philosophy. **We don't prompt our AI.** We allocate the right skills to each task, wire them into efficient workflows, and let the loops optimise themselves. Together with the second brain, it means the box gets *both* smarter (more context) and *cheaper and faster* (more efficient) every day it runs — two compounding advantages no rented, stateless cloud AI can offer.

4E. The R&D Tax Incentive — the appliance can help fund itself

Adopting Omniscient isn't just buying software — for an eligible Australian company it can be run as genuine R&D. Replacing capped subscriptions with a sovereign, always-on appliance is a real experiment in new ways of working, and that experimentation can qualify for the **R&D Tax Incentive**.

An example hypothesis: *"We hypothesise that an on-site, locally-served AI model — replacing subscription services — lets our team develop new workflows and methods that increase productivity, reduce costs, fully secure our institutional knowledge, eliminate context loss, and improve throughput through always-on AI."* Systematic, hypothesis-driven work to develop and test new methods, where the outcome isn't known in advance, is exactly what the R&DTI is designed to support.

What can be claimed — worked example, a 10-person team: - 10 people using the AI ~80% of the day to design and test new AI-enabled workflows → roughly 80% of those salaries, apportioned to the eligible activity. - Plus the appliance's **3-year lease payments, to the extent it is used in the R&D** — claimable as incurred, because you *lease* the hardware (Omniscient retains ownership) rather than own it. - For companies under **$20M turnover**, eligible R&D expenditure attracts a **43.5% refundable offset** — cash back, even from a tax-loss position.

The appliance documents its own R&D. R&DTI claims live or die on contemporaneous records — and the box's tamper-evident audit ledger already timestamps every experiment, run and result: exactly the evidence a claim needs. It's a rare alignment — the same engine that proves your privacy controls also substantiates your tax claim.

> **Not tax advice.** Eligibility must be assessed, activities registered with AusIndustry, and contemporaneous records kept; companies over $20M turnover receive a non-refundable offset at a different rate. Omniscient (with Carbon Project Australia's R&DTI practice) can help structure the activities and substantiate the claim.

5. The security story (why a CISO says yes)

Omniscient is engineered for the questions a serious security team will actually ask:

- **"Can our data leave?"** No inbound internet path, ever. A single default-deny egress broker is the only way out, and it permits only two pinned destinations: the consented burst endpoint and the signed update mirror. Flip one switch — or the physical air-gap — and the box is an island.
- **"Can the AI go rogue once it's wired into everything?"** Agents can't call tools directly. Every action runs through a policy broker with risk classes; anything consequential (money, personal data, external sends, deletions) requires human or dual approval — no matter what the model decided. Instructions hidden in incoming emails or documents are treated as untrusted and can't escalate privilege.
- **"Can Omniscient see our data? Is there a backdoor?"** No standing access. Support sessions are initiated by you, time-boxed, scope-limited, and recorded to *your* audit ledger. You hold the recovery quorum of encryption keys; our escrow share alone can't decrypt anything. You can lock us out. You can never be locked out.
- **"How do you keep it updated without a backdoor?"** The box *pulls* cryptographically signed updates — never a push, never an open port — verifies them, and applies them as atomic, auto-rolling-back image swaps inside windows you approve. Model updates, software updates, and compliance updates all arrive the same trusted way.

Everything above produces evidence. The audit ledger is hash-chained and tamper-evident: when the regulator, the board, or a customer asks what happened, you can show them.

6. Australian compliance, built in

The box is designed to support — and to *evidence* — the obligations Australian enterprises actually carry:

- **Privacy Act 1988 and its reforms** — onshore data, no offshore disclosure by default, access/retention/deletion controls, transparency for automated decisions.
- **Australian AI guardrails** — human oversight on consequential actions, record-keeping, testing, model provenance, accountable roles.
- **Sector overlays** — finance (APRA CPS 234), health records, government data classification, defence-adjacent — addressed on hardened tiers.
- **Essential Eight alignment** — patching, application control, MFA, least privilege, backups, and logging, supported out of the box.

> Australian AI and privacy law is evolving quickly. Your compliance pack is built and maintained with an Australian privacy/AI lawyer and shipped as versioned, signed updates — which is exactly why the sovereignty subscription is a genuine, ongoing service, not rent.

7. The moat — why this is hard to copy

- **Foreign hyperscalers can't follow** without abandoning their own business model — your data *is* their product.
- **Local resellers can sell hardware** but don't have the integrated agent + compliance + evidence stack.
- **Omniscient owns the IP** — the agent fleet, the hybrid-routing engine, and the contemporaneous-evidence backbone, refined across our own platforms before it ever reached your server room.

8. How it lands — the engagement

1. **Discovery** — we audit your workflows, data, software estate, and compliance obligations.
2. **Provision** — we build and harden your box to spec and load your workflows and domain agents.
3. **Install** — physical install on-site, integrated with your network and identity.
4. **Onboard & train** — your staff are trained; the box learns your organisation.
5. **Manage** — model, software, security, and compliance updates; monitoring and support — all without exfiltrating your data.

Typical time from order to productive: weeks, not the quarters a bespoke build would take.

9. Packaging & pricing model

Component	What it is	Commercial model
Omniscient Core appliance	Hardware + hardened load-out + install	**3-year lease** — Omniscient owns the hardware (so the lease + staff salaries are R&D-claimable)
Sovereignty Subscription	Model + software + **compliance** updates, security patching, support	Annual recurring
Professional Services	Custom agents, fine-tuning, deeper integrations	Project / day-rate

Tiers: *Edge* (branch/pilot) · **Business (mid-market default)** · *Business-HA* (zero-downtime) · *Enterprise* (multi-site) · *Regulated* (defence/health/gov hardening).

- Business tier ships at a **99.5%** availability target; the **HA upgrade** reaches **99.9%+**.
- Telemetry is **operational-metrics-only, opt-out** — we can support you proactively without ever seeing your content.

(Specific dollar figures live in the separate BOM & cost model.)

10. Ideal customer profile

Mid-market Australian enterprises (≈100–1,000 staff) that:

- handle sensitive or regulated data (client records, health, financial, IP, defence-adjacent);
- have hit a wall with foreign AI on cost, privacy, or guardrails;
- carry real compliance obligations and a board that asks where the data is;
- want the upside of an AI workforce without becoming an AI platform company.

11. Objection handling

- *"On-prem AI will be weaker than the big clouds."* For the vast majority of business workloads, modern open-weight models are more than capable — and for the genuinely hard task, consented burst gives you frontier capability on demand. You get both.
- *"We can't maintain a box ourselves."* You don't. Managed mode keeps it patched and current through signed, customer-approved updates — with no standing access to your data.
- *"Isn't this just a server with some software?"* The value is the integration: AI wired into every app, governed by a policy broker, and evidenced by a tamper-proof ledger. The hardware is the least interesting part.
- *"What if Omniscient disappears?"* You own the box and the keys. It keeps running. Open standards and open-source mean no lock-in to us specifically.

12. The one sentence to leave them with

> *"Everyone else rents you intelligence and keeps your data. We hand you an AI workforce that's yours to command, that never lets your data leave the building, and that proves it."*